

Attention Is All You Need

Abstract

We propose the Transformer architecture for sequence transduction.

1 Introduction

Recurrent models have been dominant. The Transformer relies entirely on attention [1].

2 Background

- Autoregressive decoding
- Encoder-decoder stacks
- Scaled dot-product attention

3 Model Architecture

```
def attention(q, k, v):  
    return softmax(q @ k.T) @ v
```